

AI Assistant Prepared

Incident	inc-9a2b66db45
Service	payment-api
Assistant status	READY
Forecast model	hosted_time_series_ready
Forecast source	fallback_ewma
Forecast latency	1245.25 ms
Forecast confidence	0.95

Prompt

Generate or refine a Splunk SPL query for incident correlation and noise reduction. Incident inc-9a2b66db45 on payment-api has root cause Upstream dependency failure, most likely payment_gateway, with secondary latency impact on /payments/authorize and /payments/capture. and signals latency_regression, upstream_api_failure, upstream_api_failure, upstream_api_failure, cpu_saturation, latency_regression, latency_regression, upstream_api_failure. Focus on grouping by service, incident_id, 5-minute windows, signal counts, host diversity, latency anomalies, and resource pressure from cpu_pct and memory_pct.

Suggested SPL

```
index=main sourcetype="agentic-ops" (service="payment-api" OR incident_id="inc-9a2b66db45") | eval status_code=tonumber(status_code), latency_ms=tonumber(latency_ms), cpu_pct=tonumber(cpu_pct), memory_pct=tonumber(memory_pct), error_type=coalesce(error_type,"null") | eventstats avg(latency_ms) as baseline_latency_avg stdev(latency_ms) as baseline_latency_std avg(cpu_pct) as baseline_cpu_avg stdev(cpu_pct) as baseline_cpu_std avg(memory_pct) as baseline_memory_avg stdev(memory_pct) as baseline_memory_std by service | eval latency_z=(latency_ms-baseline_latency_avg)/baseline_latency_std, cpu_z=(cpu_pct-baseline_cpu_avg)/baseline_cpu_std, memory_z=(memory_pct-baseline_memory_avg)/baseline_memory_std | eval signal=case(status_code>=500, "server_error", error_type="database_timeout", "database_timeout", error_type="upstream_api_failure", "dependency_failure", error_type="auth_failure", "auth_failure", error_type="deployment_regression", "deployment_regression", error_type="cpu_saturation", "cpu_saturation", error_type="memory_pressure", "memory_pressure", error_type="latency_regression" OR latency_z>2.5, "latency_anomaly", cpu_pct>=90 OR cpu_z>2.5, "cpu_pressure", memory_pct>=90 OR memory_z>2.5, "memory_pressure", true(), "noise") | where signal!="noise" | bin _time span=5m | stats count as signal_count values(signal) as signals values(host) as hosts values(endpoint) as endpoints avg(latency_ms) as avg_latency max(latency_ms) as max_latency avg(cpu_pct) as avg_cpu_pct max(cpu_pct) as max_cpu_pct avg(memory_pct) as avg_memory_pct max(memory_pct) as max_memory_pct by _time service incident_id
```

Explanation

Prepared for Splunk AI Assistant usage. Paste the prompt into AI Assistant or use the generated SPL as the starting point for correlation tuning.

Hosted time series query

```
index=main sourcetype="agentic-ops" service="payment-api" | timechart span=5m  
avg(latency_ms) as latency_ms max(status_code) as status_code | predict latency_ms  
future_timespan=3 holdback=5
```

[Open FastAPI dashboard](#)