Splunk Agentic Ops Incident Copilot

Last updated 2026-06-14T14:19:22.666+05:30

329

Total incidents

Status

APPROVED: 4, CLOSED: 17, COMPLETED: 283, EXECUTED: 5, INVESTIGATED: 11, OPEN: 7, REJECTED: 2

Severity

CRITICAL: 17, HIGH: 130, LOW: 83, MEDIUM: 99

COMMON SERVICES

Representative service groups in this demo are checkout, payment, and auth, which match the main incident flows and dashboard filters.

NOISE REDUCTION

Filters routine successful low-latency events before RCA. Signal remains when HTTP status is 500 or higher, error_type is present, or latency reaches 1000 ms.

total_events	**500**
signal_events	**127**
noise_events	**373**
noise_reduction	**74.6%**
method	**Noise is filtered when status_code is below 500, error_type is null, and latency_ms is below 1000. Remaining events are treated as investigation signals.**

CORRELATED INCIDENTS

inc-20260614084616-352496

checkout-api | Database connection pool saturation or database timeout | 17 | 5m

inc-9a2b66db45
payment-api | Upstream dependency failure | 58 | 5m

inc-02360aa6e9
payment-api | Upstream dependency failure | 58 | 5m

inc-28073388fc
payment-api | Upstream dependency failure | 58 | 5m

inc-7d0dd9f78a
checkout-api | Database connection pool saturation or database timeout | 16 | 5m

TOP ROOT CAUSES

latency regression under load	**175**
Authentication provider or token validation failure	**48**
identity provider authentication failure	**47**
Database connection pool saturation or database timeout	**39**
Upstream dependency failure	**32**
Service latency regression	**19**
payment gateway dependency failure	**15**
Service memory pressure	**14**

MCP INVESTIGATION RESULTS

Click Investigate for full detail analysis on the selected incident, including MCP tools, evidence, AI summary, and remediation actions.

inc-20260614084616-352496
checkout-api | A deployment regression in checkout-api version 2026.06.3 caused elevated latency and HTTP 502/503/504 failures across checkout-01, checkout-02, and checkout-03, with PostgreSQL connection pool saturation/timeouts as the most likely immediate bottleneck. | splunk_run_query | MCP client call: MCP evidence for checkout-api: 11 events; errors=deployment_regression; hosts=checkout-01, checkout-02, checkout-03; endpoints=/checkout; dependencies=postgres; avg_latency_ms=2653; max_latency_ms=4083; max_cpu_pct=81; max_db_pool_pct=97.

inc-9a2b66db45
payment-api | Upstream dependency failure, most likely payment_gateway, with secondary latency impact on /payments/authorize and /payments/capture. | splunk_run_query | MCP client call: MCP evidence for payment-api: 50 events; errors=latency_regression, upstream_api_failure; hosts=payment-01, payment-02; endpoints=/payments/authorize, /payments/capture; dependencies=identity_provider, none, payment_gateway, postgres; avg_latency_ms=1269; max_latency_ms=4055; max_cpu_pct=81; max_db_pool_pct=74.

inc-02360aa6e9

payment-api | Upstream dependency failure causing elevated latency on payment-api, most likely impacting /payments/authorize and /payments/capture across payment-01 and payment-02. | splunk_run_query | MCP client call: MCP evidence for payment-api: 50 events; errors=latency_regression, upstream_api_failure; hosts=payment-01, payment-02; endpoints=/payments/authorize, /payments/capture; dependencies=identity_provider, none, payment_gateway, postgres; avg_latency_ms=331; max_latency_ms=1449; max_cpu_pct=71; max_db_pool_pct=69.

inc-28073388fc

payment-api | Upstream dependency failure affecting payment-api, most likely intermittent identity_provider or payment_gateway degradation causing retries and latency spikes. | splunk_run_query | MCP client call: MCP evidence for payment-api: 50 events; errors=latency_regression, upstream_api_failure; hosts=payment-01, payment-02; endpoints=/payments/authorize, /payments/capture; dependencies=identity_provider, none, payment_gateway, postgres; avg_latency_ms=334; max_latency_ms=1493; max_cpu_pct=71; max_db_pool_pct=69.

inc-ff068e9c5d

payment-api | Primary cause appears to be an upstream API failure on payment-api's /payments/authorize path, which drove persistent latency spikes and secondary latency regression under load; the issue is localized to the authorize endpoint rather than a broad service outage. | splunk_run_query

REMEDIATION STATUS

APPROVED	**26**
OPEN	**24**
Remediation Executed	**10**
EXECUTED	**9**
TICKET CLOSED	**7**
INVESTIGATED	**6**
CLOSED	**6**
REJECTED remediation	**2**

INCIDENT TIMELINE

inc-20260614084616-352496

splunk_webhook_triaged | checkout-api | COMPLETED

inc-9a2b66db45

ticket_closed | payment-api | CLOSED

inc-9a2b66db45

remediation_executed | payment-api | EXECUTED

inc-9a2b66db45

remediation_approved | payment-api | APPROVED

inc-9a2b66db45

splunk_webhook_triaged | payment-api | COMPLETED

MCP TOOL USAGE

splunk_run_query	**492**
splunk_get_info	**8**

CONFIDENCE SCORES

inc-20260614084616-352496

checkout-api | A deployment regression in checkout-api version 2026.06.3 caused elevated latency and HTTP 502/503/504 failures across checkout-01, checkout-02, and checkout-03, with PostgreSQL connection pool saturation/timeouts as the most likely immediate bottleneck. | 0.95

inc-9a2b66db45

payment-api | Upstream dependency failure, most likely payment_gateway, with secondary latency impact on /payments/authorize and /payments/capture. | 0.85

inc-02360aa6e9

payment-api | Upstream dependency failure causing elevated latency on payment-api, most likely impacting /payments/authorize and /payments/capture across payment-01 and payment-02. | 0.75

inc-28073388fc

payment-api | Upstream dependency failure affecting payment-api, most likely intermittent identity_provider or payment_gateway degradation causing retries and latency spikes. | 0.79

inc-ff068e9c5d

payment-api | Primary cause appears to be an upstream API failure on payment-api's /payments/authorize path, which drove persistent latency spikes and secondary latency regression under load; the issue is localized to the authorize endpoint rather than a broad service outage. | 0.91

INVESTIGATION SOURCE

Deterministic fallback	**237**
Codex fallback after failure	**161**
Codex MCP agent	**82**
Unknown	**20**

MCP METRICS

query_count	**500**

success_count	**263**
failure_count	**237**
average_investigation_time_ms	**6271.26**

AI ASSISTANT USAGE

Prepared prompts and SPL suggestions for Splunk AI Assistant live query tuning.

inc-9a2b66db45

payment-api | generate_spl | Generate or refine a Splunk SPL query for incident correlation and noise reduction. Incident inc-9a2b66db45 on payment-api has root cause Upstream dependency failure, most likely payment_gateway, with secondary latency impact on /payments/authorize and /payments/capture. and signals latency_regression, upstream_api_failure, upstream_api_failure, upstream_api_failure, upstream_api_failure, cpu_saturation, latency_regression, latency_regression, upstream_api_failure. Focus on grouping by service, incident_id, 5-minute windows, signal counts, host diversity, latency anomalies, and resource pressure from cpu_pct and memory_pct. | index=main sourcetype="agentic-ops" (service="payment-api" OR incident_id="inc-9a2b66db45") | eval status_code=tonumber(status_code), latency_ms=tonumber(latency_ms), cpu_pct=tonumber(cpu_pct), memory_pct=tonumber(memory_pct), error_type=coalesce(error_type,"null") | eventstats avg(latency_ms) as baseline_latency_avg stdev(latency_ms) as baseline_latency_std avg(cpu_pct) as baseline_cpu_avg stdev(cpu_pct) as baseline_cpu_std avg(memory_pct) as baseline_memory_avg stdev(memory_pct) as baseline_memory_std by service | eval latency_z=if(baseline_latency_std>0,(latency_ms-baseline_latency_avg)/baseline_latency_std,0) | eval cpu_z=if(baseline_cpu_std>0,(cpu_pct-baseline_cpu_avg)/baseline_cpu_std,0) | eval memory_z=if(baseline_memory_std>0,(memory_pct-baseline_memory_avg)/baseline_memory_std,0) | eval signal=case(status_code>=500, "server_error", error_type="database_timeout", "database_timeout", error_type="upstream_api_failure", "dependency_failure", error_type="auth_failure", "auth_failure", error_type="deployment_regression", "deployment_regression", error_type="cpu_saturation", "cpu_saturation", error_type="memory_pressure", "memory_pressure", error_type="latency_regression" OR latency_z>2.5, "latency_anomaly", cpu_pct>=90 OR cpu_z>2.5, "cpu_pressure", memory_pct>=90 OR memory_z>2.5, "memory_pressure", true(), "noise") | where signal!="noise" | bin _time span=5m | stats count as signal_count values(signal) as signals values(host) as hosts values(endpoint) as endpoints avg(latency_ms) as avg_latency max(latency_ms) as max_latency avg(cpu_pct) as avg_cpu_pct max(cpu_pct) as max_cpu_pct avg(memory_pct) as avg_memory_pct max(memory_pct) as max_memory_pct by _time service incident_id

inc-20260614084616-352496

checkout-api | generate_spl | Generate or refine a Splunk SPL query for incident correlation and noise reduction. Incident inc-20260614084616-352496 on checkout-api has root cause Database connection pool saturation or database timeout and signals deployment_regression. Focus on grouping by service, incident_id, 5-minute windows, signal counts, host diversity, latency anomalies, and resource pressure from cpu_pct and memory_pct. | index=main sourcetype="agentic-ops" (service="checkout-api" OR incident_id="inc-20260614084616-352496") | eval status_code=tonumber(status_code), latency_ms=tonumber(latency_ms), cpu_pct=tonumber(cpu_pct), memory_pct=tonumber(memory_pct), error_type=coalesce(error_type,"null") | eventstats avg(latency_ms) as baseline_latency_avg stdev(latency_ms) as baseline_latency_std avg(cpu_pct) as baseline_cpu_avg stdev(cpu_pct) as baseline_cpu_std avg(memory_pct) as baseline_memory_avg stdev(memory_pct) as baseline_memory_std by service | eval latency_z=if(baseline_latency_std>0,(latency_ms-baseline_latency_avg)/baseline_latency_std,0) | eval cpu_z=if(baseline_cpu_std>0,(cpu_pct-baseline_cpu_avg)/baseline_cpu_std,0) | eval memory_z=if(baseline_memory_std>0,(memory_pct-baseline_memory_avg)/baseline_memory_std,0) | eval signal=case(status_code>=500, "server_error", error_type="database_timeout", "database_timeout", error_type="upstream_api_failure", "dependency_failure", error_type="auth_failure", "auth_failure", error_type="deployment_regression", "deployment_regression", error_type="cpu_saturation", "cpu_saturation", error_type="memory_pressure", "memory_pressure", error_type="latency_regression" OR latency_z>2.5, "latency_anomaly", cpu_pct>=90 OR cpu_z>2.5, "cpu_pressure", memory_pct>=90 OR memory_z>2.5, "memory_pressure", true(), "noise") | where signal!="noise" | bin _time span=5m | stats count as signal_count values(signal) as signals values(host) as hosts values(endpoint) as endpoints avg(latency_ms) as avg_latency max(latency_ms) as max_latency avg(cpu_pct) as avg_cpu_pct max(cpu_pct) as max_cpu_pct avg(memory_pct) as avg_memory_pct max(memory_pct) as max_memory_pct by _time service incident_id

inc-02360aa6e9

payment-api | generate_spl | Generate or refine a Splunk SPL query for incident correlation and noise reduction. Incident inc-02360aa6e9 on payment-api has root cause Upstream dependency failure and signals latency_regression, latency_regression, upstream_api_failure, latency_regression, latency_regression. Focus on grouping by service, incident_id, 5-minute windows, signal counts, host diversity, latency anomalies, and resource pressure from cpu_pct and memory_pct. | index=main sourcetype="agentic-ops" (service="payment-api" OR incident_id="inc-02360aa6e9") | eval status_code=tonumber(status_code), latency_ms=tonumber(latency_ms), cpu_pct=tonumber(cpu_pct), memory_pct=tonumber(memory_pct), error_type=coalesce(error_type,"null") | eventstats avg(latency_ms) as baseline_latency_avg stdev(latency_ms) as baseline_latency_std avg(cpu_pct) as baseline_cpu_avg stdev(cpu_pct) as baseline_cpu_std avg(memory_pct) as baseline_memory_avg stdev(memory_pct) as baseline_memory_std by service | eval latency_z=if(baseline_latency_std>0,(latency_ms-baseline_latency_avg)/baseline_latency_std,0) | eval cpu_z=if(baseline_cpu_std>0,(cpu_pct-baseline_cpu_avg)/baseline_cpu_std,0) | eval memory_z=if(baseline_memory_std>0,(memory_pct-baseline_memory_avg)/baseline_memory_std,0) | eval signal=case(status_code>=500, "server_error", error_type="database_timeout", "database_timeout", error_type="upstream_api_failure", "dependency_failure", error_type="auth_failure", "auth_failure", error_type="deployment_regression", "deployment_regression", error_type="cpu_saturation", "cpu_saturation", error_type="memory_pressure", "memory_pressure", error_type="latency_regression" OR latency_z>2.5, "latency_anomaly", cpu_pct>=90 OR cpu_z>2.5, "cpu_pressure", memory_pct>=90 OR memory_z>2.5, "memory_pressure", true(), "noise") | where signal!="noise" | bin _time span=5m | stats count as signal_count values(signal) as signals values(host) as hosts values(endpoint) as endpoints avg(latency_ms) as avg_latency max(latency_ms) as max_latency avg(cpu_pct) as avg_cpu_pct max(cpu_pct) as max_cpu_pct avg(memory_pct) as avg_memory_pct max(memory_pct) as max_memory_pct by _time service incident_id

inc-28073388fc

payment-api | generate_spl | Generate or refine a Splunk SPL query for incident correlation and noise reduction. Incident inc-28073388fc on payment-api has root cause Upstream dependency failure affecting payment-api, most likely intermittent identity_provider or payment_gateway degradation causing retries and latency spikes. and signals latency_regression. Focus on grouping by service, incident_id, 5-minute windows, signal counts, host diversity, latency anomalies, and resource pressure from cpu_pct and memory_pct. | index=main sourcetype="agentic-ops" (service="payment-api" OR incident_id="inc-28073388fc") | eval status_code=tonumber(status_code), latency_ms=tonumber(latency_ms), cpu_pct=tonumber(cpu_pct), memory_pct=tonumber(memory_pct), error_type=coalesce(error_type,"null") | eventstats avg(latency_ms) as baseline_latency_avg stdev(latency_ms) as baseline_latency_std avg(cpu_pct) as baseline_cpu_avg stdev(cpu_pct) as baseline_cpu_std avg(memory_pct) as baseline_memory_avg stdev(memory_pct) as baseline_memory_std by service | eval latency_z=if(baseline_latency_std>0,(latency_ms-baseline_latency_avg)/baseline_latency_std,0) | eval cpu_z=if(baseline_cpu_std>0,(cpu_pct-baseline_cpu_avg)/baseline_cpu_std,0) | eval memory_z=if(baseline_memory_std>0,(memory_pct-baseline_memory_avg)/baseline_memory_std,0) | eval signal=case(status_code>=500, "server_error", error_type="database_timeout", "database_timeout", error_type="upstream_api_failure", "dependency_failure", error_type="auth_failure", "auth_failure", error_type="deployment_regression", "deployment_regression", error_type="cpu_saturation", "cpu_saturation", error_type="memory_pressure", "memory_pressure", error_type="latency_regression" OR latency_z>2.5, "latency_anomaly", cpu_pct>=90 OR cpu_z>2.5, "cpu_pressure", memory_pct>=90 OR memory_z>2.5, "memory_pressure", true(), "noise") | where signal!="noise" | bin _time span=5m | stats count as signal_count values(signal) as signals values(host) as hosts values(endpoint) as endpoints avg(latency_ms) as avg_latency max(latency_ms) as max_latency avg(cpu_pct) as avg_cpu_pct max(cpu_pct) as max_cpu_pct avg(memory_pct) as avg_memory_pct max(memory_pct) as max_memory_pct by _time service incident_id

inc-7d0dd9f78a

checkout-api | generate_spl | Generate or refine a Splunk SPL query for incident correlation and noise reduction. Incident inc-7d0dd9f78a on checkout-api has root cause Database connection pool saturation or database timeout and signals database_timeout. Focus on grouping by service, incident_id, 5-minute windows, signal counts, host diversity, latency anomalies, and resource pressure from cpu_pct and memory_pct. | index=main sourcetype="agentic-ops" (service="checkout-api" OR incident_id="inc-7d0dd9f78a") | eval status_code=tonumber(status_code), latency_ms=tonumber(latency_ms), cpu_pct=tonumber(cpu_pct), memory_pct=tonumber(memory_pct), error_type=coalesce(error_type,"null") | eventstats avg(latency_ms) as baseline_latency_avg stdev(latency_ms) as baseline_latency_std avg(cpu_pct) as baseline_cpu_avg stdev(cpu_pct) as baseline_cpu_std avg(memory_pct) as baseline_memory_avg stdev(memory_pct) as baseline_memory_std by service | eval latency_z=if(baseline_latency_std>0,(latency_ms-baseline_latency_avg)/baseline_latency_std,0) | eval cpu_z=if(baseline_cpu_std>0,(cpu_pct-baseline_cpu_avg)/baseline_cpu_std,0) | eval memory_z=if(baseline_memory_std>0,(memory_pct-baseline_memory_avg)/baseline_memory_std,0) | eval signal=case(status_code>=500, "server_error", error_type="database_timeout", "database_timeout", error_type="upstream_api_failure", "dependency_failure", error_type="auth_failure", "auth_failure", error_type="deployment_regression", "deployment_regression", error_type="cpu_saturation", "cpu_saturation", error_type="memory_pressure", "memory_pressure", error_type="latency_regression" OR latency_z>2.5, "latency_anomaly", cpu_pct>=90 OR cpu_z>2.5, "cpu_pressure", memory_pct>=90 OR memory_z>2.5, "memory_pressure", true(), "noise") | where signal!="noise" | bin _time span=5m | stats count as signal_count values(signal) as signals values(host) as hosts values(endpoint) as endpoints avg(latency_ms) as avg_latency max(latency_ms) as max_latency avg(cpu_pct) as avg_cpu_pct max(cpu_pct) as max_cpu_pct avg(memory_pct) as avg_memory_pct max(memory_pct) as max_memory_pct by _time service incident_id

HOSTED TIME SERIES FORECAST

Forecast-ready signals for a hosted time-series model or fallback forecast pipeline.

1
Predicted Critical Incidents

3
Services At Risk

92%
Avg Forecast Confidence

SERVICE	RISK	ETA	CONFIDENCE
payment-api	Healthy	-	95%
checkout-api	Critical	15 min	85%
payment-api	Healthy	-	95%

Incident actions for inc-20260614084616-352496 / checkout-api

Investigate
Prepare SPL
Approve
Execute
Reject
Close

Showing 1-5 of 329 incidents | Page 1 of 66 | Previous | Next

INCIDENT	SERVICE	SEVERITY	REMEDIATION	ROOT CAUSE	MCP EVIDENCE	AI SUMMARY	ACTIONS
inc-20260614084616-352496	checkout-api	**CRITICAL**	**INVESTIGATED**	A deployment regression in checkout-api version 2026.06.3 caused elevated latency and HTTP 502/503/504 failures across. **More** codex_mcp_agent; MCP=True; events=11	MCP client call: MCP evidence for checkout-api: 11 events; errors=deployment_regression; hosts=checkout-01,. **More**	AI assistant workflow was not used. The RCA was derived from the Splunk query result and the supplied incident context, which independently point to a deployment regression with PostgreSQL connection pool pressure as the proximate cause. Reasoning: this summary interprets the evidence rather than repeating it. MCP supplied 11 correlated events; the RCA confidence is 0.95, so the recommended operator path is: Recycle saturated DB connections; Scale checkout-api workers; Escalate to database owner.	Investigate Approve Execute Reject Close
inc-9a2b66db45	payment-api	**HIGH**	**CLOSED**	Upstream dependency failure, most likely payment_gateway, with secondary latency impact on /payments/authorize and. **More** codex;	MCP client call: MCP evidence for payment-api: 50 events; errors=latency_regression, upstream_api_failure;. **More**	Telemetry supports an upstream dependency outage or degradation as the primary driver of the incident. The service shows elevated latency and repeated upstream_api_failure errors while host CPU, memory, and DB pool usage remain below	Investigate Approve Execute Reject Close

INCIDENT	SERVICE	SEVERITY	REMEDIATION	ROOT CAUSE	MCP EVIDENCE	AI SUMMARY	ACTIONS
				MCP=True; events=50		saturation, making an internal capacity issue less likely. The incident is consistent with dependency-sensitive request paths failing or retrying excessively. Reasoning: this summary interprets the evidence rather than repeating it. MCP supplied 50 correlated events; the RCA confidence is 0.85, so the recommended operator path is: Fail over or route around the degraded upstream dependency if an alternate is available.; Open a provider escalation for payment_gateway and verify identity_provider health if authorization paths are affected.; Throttle or cap retries to reduce amplification while the dependency is unstable..	
inc-28073388fc	payment-api	**HIGH**	**CLOSED**	Upstream dependency failure affecting payment-api, most likely intermittent identity_provider or payment_gateway degradation. **More** codex; MCP=True; events=50	MCP client call: MCP evidence for payment-api: 50 events; errors=latency_regression, upstream_api_failure;. **More**	Pattern is consistent with upstream dependency instability rather than a local service outage. Both hosts and both payment endpoints were impacted, while CPU, memory, and DB pool remained below critical limits. The incident was completed using fallback rules after AI was unavailable, so the RCA confidence is moderate and should be validated against upstream provider health and retry/error-rate	Investigate Approve Execute Reject Close

INCIDENT	SERVICE	SEVERITY	REMEDIATION	ROOT CAUSE	MCP EVIDENCE	AI SUMMARY	ACTIONS
						trends. Reasoning: this summary interprets the evidence rather than repeating it. MCP supplied 50 correlated events; the RCA confidence is 0.79, so the recommended operator path is: Fail over to a healthy upstream dependency path; Open escalation with the upstream provider team; Throttle or cap retries to reduce amplification.	
inc-02360aa6e9	payment-api	**HIGH**	**INVESTIGATED**	Upstream dependency failure causing elevated latency on payment-api, most likely impacting /payments/authorize and. **More** codex; MCP=True; events=50	MCP client call: MCP evidence for payment-api: 50 events; errors=latency_regression, upstream_api_failure;. **More**	The incident is best explained by degraded upstream dependency behavior rather than local resource exhaustion. Latency spikes are widespread across both payment hosts and both payment endpoints, while CPU and DB pool usage are elevated but not at failure thresholds. The pattern supports a dependency-driven outage or partial degradation affecting request paths in payment-api. Reasoning: this summary interprets the evidence rather than repeating it. MCP supplied 50 correlated events; the RCA confidence is 0.75, so the recommended operator path is: Fail over or reroute traffic away from the degraded upstream dependency.; Open escalation with the upstream provider or owning team	Investigate Approve Execute Reject Close

INCIDENT	SERVICE	SEVERITY	REMEDIATION	ROOT CAUSE	MCP EVIDENCE	AI SUMMARY	ACTIONS
						immediately.; Throttle retries and add backoff to prevent retry amplification..	
inc-ff068e9c5d	payment-api	**LOW**	**INVESTIGATED**	Primary cause appears to be an upstream API failure on payment-api's /payments/authorize path, which drove persistent latency. **More** codex_mcp_agent; MCP=True; events=342	No MCP evidence yet	The evidence points to a localized dependency failure on payment-api /payments/authorize. Latency regression is a downstream symptom; upstream_api_failure is the stronger primary signal. Reasoning: this summary interprets the evidence rather than repeating it. MCP supplied 342 correlated events; the RCA confidence is 0.91, so the recommended operator path is: Inspect the upstream dependency behind /payments/authorize for errors, timeouts, or degraded responses.; Review retry, timeout, and circuit-breaker settings on payment-api for the authorize flow.; Scale payment-api instances if elevated latency and CPU pressure persist..	Investigate Approve Execute Reject Close